Overview

This project is meant to evaluate the quality of translation and localization of various LLM models when asked to translate emails from English to a given language and locale. Evaluate the email from the perspective of a person in the targeted country who speaks the language and has just received this message in their inbox without knowing it was AI-generated. How well is it translated? How well does it target local traditions and norms? How true to the original content and tone is the translation?

We'll ask several questions, all using the same basic evaluation scale. Keep these ratings and descriptions in mind while you are evaluating.

- did not translate – select this if the model utilises the original English instead of translating the text
- serious failures exist - use this in cases where you would be very disappointed, confused, offended, or in some other way have negative feelings towards the company or product because of the content of the translation
- imperfect but not terrible - there are errors or issues that are very noticeable, but that are not so bad as to give a negative impression of the company, the main ideas come through and it is clear what is being advertised.
- mostly good with small issues - the wording or translations are noticably non-native, or are awkward or a little odd, but it is overall something that makes sense and could be used without embarassment on the part of the company.
- very good or nearly perfect - this is for something that seems very close to natural, native, and culturally appropriate.

Steps

1. Read the original English email and the translation.

2. Answer segment specific questions, considering the context and intended audience of the email.

3. Leave free-text comments to add context to your rating.

Notes

- The translated email may include notes from the model on the translation. Please disregard these and evaluate the translated email from the perspective of a potential customer receiving it in your inbox.